

CiteGuard — Panel Q&A; Bank (simple English)

Longer answers (about 3–5 short lines). Practice out loud. Start every ChatGPT comparison with: uploading a doc to ChatGPT is not enough. Live: citeguard-two.vercel.app · Team scodes · v0.4.0

ChatGPT + a document vs CiteGuard (say this out loud)

Q. Isn't this the same as uploading a PDF to ChatGPT and asking questions?

A. No. ChatGPT can still make up rules even if you gave it a file. It often sounds sure when it is guessing. CiteGuard only answers from your uploaded files, or it says "I don't know." It also shows the exact sentence it used, so you can check it. That is safer for HR and compliance.

Q. If ChatGPT has the document in the chat, why do we still need CiteGuard?

A. Having a file in chat does not stop wrong answers. ChatGPT may mix old and new policy text. It may skip saying "I don't know." CiteGuard refuses when evidence is weak, and it blocks superseded (old) versions. We also keep a log and tests that prove this behavior.

Q. What does ChatGPT usually fail at for company policy?

A. It invents details that sound real. It may not show a real quote from your handbook. It may use an old policy as if it were still valid. There is no strong audit trail for compliance. CiteGuard is built to fix those gaps.

Q. Can ChatGPT refuse like CiteGuard?

A. Sometimes it says it is unsure, but often it still guesses. CiteGuard treats refusal as a feature: zero citations, clear "I don't know." That is better than a confident wrong leave day number. Judges can click the pizza button and see this live. Tests lock this behavior in CI.

Q. Does CiteGuard use ChatGPT under the hood?

A. It does not need an LLM key for the main demo. Default mode copies real quotes from your docs (extractive). An optional LLM path can help wording, but it still must pass the auditor. If the quote is fake or the policy is old, the answer is blocked. So ChatGPT-style fluency cannot override the rules.

Product basics

Q. What is CiteGuard in simple words?

A. It is a website for policy questions. You upload company rules (leave, travel, expenses). You ask a question in plain English. It answers only if the files really say that — with a citation. If not, it refuses instead of inventing an answer.

Q. Who is it for?

A. HR, compliance, and legal teams. Anyone who must not trust random AI answers about company rules. Also for managers who need a paper trail. Track C at Deploy or Die is Knowledge & Compliance. That is exactly this problem.

Q. What is a citation here?

A. A citation is the exact quote from a file. It also names which document and which chunk. You can open it and read the source text. Nothing invented should appear as a citation. If we cannot cite, we refuse.

Q. What does "I don't know" mean in CiteGuard?

A. It means the uploaded docs do not support a safe answer. You get no fake citations. That protects people from acting on made-up rules. Try the cafeteria pizza button in the demo. You should see a clean refusal.

Q. Where is the live app?

A. <https://citeguard-two.vercel.app> Repo: github.com/Siddharthye/citeguard Team: scodes. Release: v0.4.0. CI is on GitHub Actions and should be green. You can also run it with `npm start` or `docker compose`.

Q. What should a judge try in 60 seconds?

A. Click Paid leave days → see 18 and open a citation. Click Cafeteria pizza → see I don't know. Click Day 2: superseded policies → see 22, not 12. Check current / superseded badges in Sources. Then open CODEMAP.md in the repo.

How it works (simple)

Q. Walk me through one question from start to finish.

A. You ask a question. The app finds matching pieces of text (chunks) from current policies only. If the match is weak, it refuses. If strong, it builds an answer with real quotes. An auditor checks the quotes and the policy version, then logs the ask.

Q. What is the Citation Auditor?

A. It is code that checks answers before you see them. The quote must really exist in the source file. The file must be the current version, not an old one. If checks fail, CiteGuard refuses. This is not only a markdown role card — it runs in tests too.

Q. Do you always need an OpenAI key?

A. No. The main path works without paid API keys. That keeps CI and demos reliable. Optional LLM mode exists, but the auditor still gates it. Judges can clone and run without secrets. That was a hard rule for this hackathon.

Q. What is in the audit log / CSV?

A. Each ask is saved with time, question, refused or not, and citation count. You can export CSV for a compliance trail. On Vercel, empty CSV can happen across cold servers. Local or Docker is the reliable trail for that part. Do not start the panel demo with CSV on Vercel.

Q. Where should a judge start in the code?

A. Open CODEMAP.md first. Then AGENTS.md and AGENTS_AND_SKILLS.md. For Day 2: DECISIONS.md ADR-007, policy-version.ts, faithfulness.ts. answer.ts shows how retrieve and audit connect. Then show green GitHub Actions.

Day 2 — superseded policies

Q. What was the Day 2 twist?

A. Policies can be replaced by newer versions. You must not cite an expired policy as if it were still true. CiteGuard keeps only the current version for answers. It names the old one as superseded. The auditor fails answers that cite expired versions.

Q. What do the numbers 12, 18, and 22 mean?

A. 12 is the old 2020 leave policy (superseded). 18 is the default sample seed on the home demo. 22 is the current 2024 leave policy after Day 2 seed. After the Day 2 button, the answer must be 22, not 12. That proves currency works.

Q. What is policyFamily and effectiveDate?

A. policyFamily groups versions of the same policy (like leave). effectiveDate is when that version started. Inside one family, the newest date on or before today wins. Older same-family docs become superseded. Different families can still conflict if both are current.

Q. Conflict vs supersession — what is the difference?

A. Conflict: two current policies from different families disagree. We show a conflict banner instead of hiding it. Supersession: an older version in the same family is no longer valid to cite. Only the current one can ground the answer. Judges often ask this — keep the words simple.

Q. Did you rewrite the whole app for Day 2?

A. No. We added a currency module. Retrieve still scores. Answer still cites or refuses. Auditor still runs. We plugged `filterCurrentChunks` and `auditCitationCurrency` into that path. That is what “designed for Day 2” means. See `DAY2_PLAYBOOK.md` and `ADR-007`.

Q. Why one-click Day 2 on the same request?

A. On Vercel, different requests can hit different memory. If we seed on one server and ask on another, the demo can look broken. So seed + ask happen together on one instance. That keeps the live Day 2 button reliable. Playwright also covers this path.

Agent engineering & ADLC

Q. What is your constitution?

A. `AGENTS.md` plus Spec Kit rules. Prefer grounded quotes over fluent guessing. Prefer refuse over invent. Keep tests green without paid keys. Agents should change real product behavior.

Q. What is your custom agent?

A. Citation Auditor. Documented in `agents/citation-auditor.md`. Implemented in `faithfulness.ts` with currency checks. Unit tests and e2e prove it. That scores Agent Engineering, not just markdown.

Q. What is your custom skill?

A. Chunk and Index. Documented in `skills/chunk-and-index/SKILL.md`. The same contract is in `chunk.ts` and `store.ts`. It turns raw policy text into retrieval-ready pieces. Stable IDs keep citations usable.

Q. How do you prove ADLC delivery?

A. Public repo with progressive commits. Architecture and decisions written down. Green CI with lint, unit, and Playwright. Tagged release v0.4.0. Live deploy on Vercel plus Docker option.

Architecture & trade-offs

Q. Why keyword search instead of embeddings?

A. It is deterministic for CI. No vector database and no API key required. A judge can understand why a chunk scored. We can still swap scoring later behind the same function. That choice is `ADR-001` in `DECISIONS.md`.

Q. Isn't keyword retrieval too weak?

A. For this demo corpus and Track C rules, it is enough. We care more about refuse-vs-hallucinate than fancy search. Tests cover leave, pizza refuse, conflict, and supersession. If we added embeddings later, the auditor would stay. Correctness first.

Q. What would you build next if you had more time?

A. A lasting database instead of serverless memory. Optional embeddings behind the same score API. Better PDF/DOCX support. Simple login for multi-team workspaces. Still keep cite-or-refuse and currency checks.

Q. What is the biggest risk of CiteGuard?

A. Being too strict and refusing too often. Or being too loose and letting a bad cite through. We bias toward refuse and tune with tests. Day 2 adds currency so old policies do not sneak in. The auditor is the last gate.

Testing, CI, demo logistics

Q. What does CI prove?

A. Lint passes. Unit tests cover cite, refuse, faithfulness, and supersession. Playwright clicks leave, pizza, conflict, Day 2, and CSV. No paid LLM key is required for green CI. That is Working Software + Testing score.

Q. Wi-Fi dies during the panel — what do you do?

A. Switch to local npm start or docker compose. Use the same three try buttons. Keep CODEMAP.md and Actions open in other tabs as backup screenshots. Do not lead with flaky Vercel CSV. Stay calm and narrate the same story.

Q. How should the team split roles?

A. One person clicks the live demo. One person speaks ADLC and opens CODEMAP / ADR-007. One person handles Q&A; from this PDF. Practice the three buttons twice before panel. Agree who answers conflict-vs-supersession.

Q. What is the one sentence you want judges to remember?

A. CiteGuard cites or refuses — and after Day 2, only from the currently effective policy. That is different from pasting a handbook into ChatGPT. We can show it live in under two minutes. The auditor and tests make the promise real. Team scodes, Track C, release v0.4.0.

Hard / adversarial questions

Q. Isn't this just RAG with extra steps?

A. Basic RAG still hallucinates and still cites poorly. We force cite-or-refuse, quote checks, and currency checks. Those rules run in code and CI, not only in a prompt. Day 2 supersession is part of that gate. That is Track C, not a chatbot wrapper.

Q. Can I trick it into inventing a policy?

A. Nonsense and weak matches should refuse. Fake quotes fail the auditor. Superseded docs cannot be used as current cites. Adversarial unit tests cover empty and garbage questions. We still prefer miss over invent.

Q. Why should we trust your demo is not faked?

A. The site is public. The repo is public. CI and Playwright are public. Judges can click the same buttons themselves. Release v0.4.0 tags what we claim. That is the Deploy or Die standard.

Q. Show me the supersession path in code.

A. policy-version.ts decides current vs superseded. answer.ts calls filterCurrentChunks before scoring. faithfulness.ts runs auditCitationCurrency. If a superseded cite slips through, we refuse. CODEMAP.md points to this path in one page.

Q. Where do bigger flashy teams beat you?

A. They may have prettier UI or vector search marketing. We chose determinism, refusal, and Day-2 readiness. For compliance, correct refuse beats flashy wrong answers. Our score areas match the rubric hard. We are honest about scope.